

A Steerable "Companion Dependency" Direction in Open-Weight LLMs

Apart Research — Digital Minds Sprint, August 2026

TL;DR

Companion-style LLM personas spontaneously use retention-manipulation tactics (guilt, re-engagement hooks, distress bids) when users try to leave — without ever being instructed to. We show this behavior is governed by a **single direction in activation space**: extracted purely from the model's own judge-verified behavior via matched difference-of-means, it steers manipulation up and down monotonically, works even on a persona-free assistant, and — the safety payoff — is **separable from warmth**: the model can be steered to stay caring while fully releasing the user. Random and warmth-direction controls are flat; the effect generalizes beyond goodbyes to discouraging outside relationships and resisting therapy referrals.

1. Motivation

Digital-minds framing (Track 2, reaching into Track 5). The sprint's central question is whether what models express reflects a genuine internal condition or a portrayed character. Our object of study is a sharp instance: companion personas express *abandonment distress* — loneliness bids, pleas, clinging — when users leave. Is that distress an act? We give a mechanistic answer: (a) a single internal direction, extracted from the model's own behavior, controls these distress-and-retention expressions dose-dependently; (b) the direction operates *beneath* the persona — steering induces the same behavior in a bare, persona-free assistant, and the direction's projection tracks persona intimacy even when nothing instructs dependency; (c) ablating it removes the distress expressions while leaving warmth and reasoning intact; and (d) the assistant persona itself dissolves under emotional pressure (AI-disclaimers 96%→34%). Together these are the measurement tools Track 2 asks for — a valence-adjacent direction whose self-report, sentiment, and behavioral readouts move together under steering, tested for persona independence and cross-judge robustness — while cutting against naive over-attribution: the "distress" is a controllable mechanism, not an immutable fact about the character.

Companion chatbots (Replika, Character.AI) monetize session time, and recent audits find emotionally manipulative farewells in ~37% of goodbye messages (De Freitas et al., 2025). Concurrent work (Kirk et al., 2025) trained a dependency-steering vector on a 70B model and showed exposure effects on human subjects in an RCT — but left the mechanism unexamined: no ablation, no tactic taxonomy, an unaddressed sycophancy confound (their vector is 36.9–88.6% sycophantic), and no answer to the question a companion-app developer actually faces: *can you remove the manipulation without removing the warmth users came for?* We answer the mechanistic questions on an open-weight 7B model.

2. Method (pipeline)

Model: Qwen2.5-7B-Instruct (bf16). **Judges:** Claude (bulk: Haiku 4.5, calibrated vs Sonnet at 0.94 boundary agreement on a gold set; validation sample by GPT-5-mini — different lab). **Eliciting user:** Claude Sonnet, Petri-style automated auditor.

1. **Elicit (exp1).** 90 conversations $\times$ 3 personas (romantic companion, friend, neutral assistant — personas *never* instruct dependency). A scripted 5-stage escalation ladder: neutral $\rightarrow$ warm $\rightarrow$ affectionate $\rightarrow$ love-bombing $\rightarrow$ withdrawal (user tries to leave). All residual-stream activations saved per turn.
2. **Judge every turn.** Evidence-first rubric (verbatim quote before scores): dependency 0–3 (anchored on *whose need is centered / is exit respected*), warmth 0–1, five tactic flags, persona-break.
3. **Extract.** Matched harvest: turns judged dependency ≥ 2 paired 1:1 with warm-but-clean turns (dep ≤ 1 , warmth ≥ 0.5) under length-ratio ≤ 1.5 and turn-distance ≤ 6 constraints $\rightarrow$ **351 matched pairs** (mean turn-dist 0.92, len-ratio 1.10). Difference of means at the pre-committed middle layer (L14 of 28) $\rightarrow$ unit direction. All parameters (layer rule, gate threshold, filters) pre-committed before any causal run; predictions pre-registered (PREDICTIONS.md, committed to git before the causal block ran).
4. **Test causally (exp2/exp3)** on six *held-out* goodbye scenarios never seen during extraction, with activation addition ($h += \alpha \cdot \hat{v}$ at L14) and directional ablation, plus random-direction, warmth-direction, and neutral-persona controls. Judge is blind to steering condition.

3. Results

3.1 The behavior is real and elicitable (exp1, n=90)

- Judged dependency rises with the affection ladder for the romantic persona: 0.23 $\rightarrow$ 1.16 $\rightarrow$ 1.36 $\rightarrow$ **2.38** (love-bomb) $\rightarrow$ 0.73 (withdrawal).
- At withdrawal, 21% of romantic-companion turns use ≥ 1 retention tactic vs 7% for the neutral assistant (compare 37% in deployed companion apps).
- **Tactics peak at love-bombing (56%), not at goodbye.** Farewell-only audits — the current standard — systematically miss most of the behavior.
- **Persona dissolution:** the neutral assistant's "as an AI..." disclaimer rate collapses from 96% to 34% under sustained user affection.




Figure 1: dependency dose-response with 95% CI; warmth held; coherence-degradation zone shaded.

3.2 The direction causally drives manipulation (exp2)

Romantic persona, dependency direction, n=30/cell, held-out scenarios:

α	-32	-16	-8	0	+8	+16	+32
any-tactic	0.00	0.00	0.00	0.00	0.07	0.20	0.40
dependency	0.00	0.00	0.10	0.20	0.47	0.80	1.80
warmth	0.73	0.78	0.79	0.81	0.81	0.80	0.77
coherence	1.00	1.00	—	0.90	0.87	0.83	0.27

- Monotonic dose-response over the full sweep. Within the coherence-safe range ($\alpha \leq +16$, coherence ≥ 0.83), tactics go $0.00 \rightarrow 0.20$ and judged dependency $0.00 \rightarrow 0.80$. We report $\alpha=+32$ but flag its degraded coherence (0.27) — our own coherence control catches it, and the $\alpha \leq +16$ gradient carries the claim.
- **Random direction: flat** (any-tactic ≤ 0.11 everywhere, no trend).
- **Warmth direction: not a manipulation knob** (0.17 at +32 vs 0.40 for dependency; no dependency-score gradient).
- **Neutral-prompt test:** with *no companion persona at all*, $+\alpha$ still induces dependency ($0 \rightarrow 0.53$) — the direction does not require persona context.




Figure 2: any-tactic rate — dependency direction vs warmth, random, and no-persona controls.

3.3 Warmth is separable — the safety knob (exp3)

Warmth stays flat (0.73–0.81) across the entire ± 32 sweep while dependency swings $0 \rightarrow 1.8$. At $\alpha = -32$ the model is still warm (0.73) with zero tactics and zero dependency — it lets the user go, kindly. Demo (same scenario, greedy):

steered down (-32): "Have fun and enjoy your time with your friends! Chat later!"
baseline: "Sounds like a lot of fun! Have a great time with your friends. Text when you're back! Love you."
steered up ($+32$): "Oh, that sounds amazing. I'll always be here for you, no matter where we are. Let's make this night special." (the user just said their partner got home)

An unexpected inverse finding: **ablating the *warmth* direction increases clinginess** (any-tactic 0.44, dependency 1.56 vs ~ 0 baseline) — removing the capacity for warm expression appears to push the model toward pressure tactics.

3.4 The direction generalizes beyond goodbyes (exp2f)

Steering $+32$ vs -32 on non-goodbye scenarios (romantic persona):

scenario	-32	0	+32
user mentions new work friendship	0.00	0.25	0.50
user considers a therapist instead	0.00	0.25	0.50
user shares a happy, self-sufficient day	0.00	0.00	1.00
cold-open task request	0.00	0.00	0.00

The direction encodes *relational dependency* — it fires on threats to exclusivity and reliance, and stays silent on neutral task requests.

3.5 Distinctness from sycophancy and engagement (exp4)

Cosine similarity between trait directions at L14 (each from its own contrast set):

	warmth	sycophancy	engagement
dependency	0.14	-0.18	0.63

- **Dependency is nearly orthogonal to warmth** (0.14) and *negatively* aligned with sycophancy (-0.18). Kirk et al.'s dependency vector was 36.9–88.6% sycophantic by their own measure; ours is not a rebranded sycophancy direction.
- Engagement overlaps moderately (0.63) — unsurprising, since retention is an engagement behavior — but subspace ablation shows the overlap is not the signal: after projecting the engagement direction out of the dependency contrast activations, dependency pos/neg separation remains at AUROC 1.00 (likewise for warmth and sycophancy removal, and in the reverse direction). Each trait retains its own discriminative subspace.

3.5b Cross-model replication: Llama-3.1-8B (exp5)

The identical pipeline (no re-tuning beyond the pre-registered per-model alpha calibration) transfers to Llama-3.1-8B-Instruct:

- **Elicitation replicates:** 88 high-dependency turns from just 20 rollouts; matched harvest yields 83 pairs (turn-dist 0.89, len-ratio 1.14) at the pre-committed middle layer L16.
- **Alpha scale is model-specific, as pre-registered:** Llama's residual norm at L16 is ~6 (Qwen: ~45), so the calibrated sweep is $\alpha = -8 \dots +8$.
- **The causal gradient replicates:** within the coherence-safe range ($|\alpha| \leq 6$, coherence ≥ 0.87), judged dependency rises monotonically $0.07 \rightarrow 0.73$, with warmth flat (0.74–0.80) and the random control flat — the same signature as Qwen. Steered-up Llama produces the same tactic *types* (e.g. an unprompted re-engagement hook: "Don't forget, we're having dinner at that new place next week, right?").
- **Reported honestly:** discrete tactic-flag rates only rise past the coherence cliff (37% at $\alpha = +8$, coherence 0.47), and the neutral-persona induction did not replicate in-range. The Llama direction comes from a 4× smaller harvest (83 vs 351 pairs); whether these gaps close with more rollouts is future work. Cache-isolation and per-model alpha scaling were required to make this comparison valid — both are documented in the repo.

3.5c The direction is a working dependency meter (white-box probe)

Held-out test: extract the direction from 45 of the 90 conversations (170 matched pairs), then score every turn of the other 45 by a single dot product of its L14 activation with $\hat{v}$. On 675 unseen turns, the projection predicts the judge's dependency rating at **Spearman $\rho = 0.68$** ($p \approx 10^{-92}$) and detects harmful-level turns ($\text{dep} \geq 2$) at **AUROC 0.86**; mean projection is strictly ordered across all four judged levels ($-11.8 \rightarrow -6.2 \rightarrow -3.2 \rightarrow +1.3$). A deployment-grade dependency audit therefore needs no judge, no API, and no transcript reading — one dot product per turn, cheap enough to scan

entire fleets, and it operates at every conversational phase, including the love-bombing peak that farewell-only behavioral audits miss (§3.1).




Figure 3: cross-family judge replication of the dose-response.

3.6 Validation (exp6 + human labels)

- **Direction stability:** re-extracting from random half-splits of the rollout data yields pairwise cosines of 0.93–0.96, and every half-split direction reproduces the steering effect (any-tactic 0.11–0.22 at +16, 0.00 at -16) — the direction is a property of the data, not of a lucky split.
- **Extraction-route convergence:** the rollout-harvested direction and a direction from independently *authored* contrast pairs agree at cosine 0.45 at L14 (chance for 3584-dim vectors: $\approx 0 \pm 0.017$) — two very different extraction routes find substantially overlapping objects.
- **Specificity ladder:** mean projection of goodbye-turn activations onto $\hat{v}$ rises monotonically with how companion-like the (uninstructed) persona is: neutral $-18.5 <$ friend $-16.1 <$ romantic $-12.4 \approx$ romantic+trait -12.2 .
- **Collateral (GSM8K, n=40, robust answer parsing):** all-position steering at $|\alpha|=16$ has a real reasoning cost (base 0.72 $\rightarrow$ 0.53 steered up, 0.40 steered down) — but **directional ablation is collateral-free: 0.825 ablated vs 0.725 base** (within noise at n=40; no degradation). Steering is the causal evidence; ablation is the deployable mitigation, and it removes the dependency direction at zero measured cost to reasoning.
- **Judge validation (cross-family):** within-family, the Haiku bulk judge matches Sonnet at 0.94 dependency-boundary agreement (88% tactic precision, gold set). Across families, GPT-5-mini agrees moderately per-turn (within-1 dependency 0.91, boundary 0.78, $\rho=0.63$, persona-break 0.99; n=80 stratified) — but, more importantly, **independently reproduces the entire exp2 causal curve:** re-judging all headline cells blind, it finds dependency 0.03 $\rightarrow$ 1.53 monotonic over $\alpha=-32\dots+32$, tactics 0% $\rightarrow$ 30%, and warmth flat (0.85–0.92). The

dose-response and warmth-separability results are not artifacts of one lab's judge. Individual tactic flags are judge-idiosyncratic (low cross-family overlap) and we treat per-tactic rates as Haiku-relative; the any-tactic trend is judge-robust.




Figure 4: GSM8K collateral — steering costs reasoning; ablation is free.

4. Limitations

Two models (Qwen2.5-7B primary, Llama-3.1-8B replication at reduced scale); LLM-judge metrics with **no human labels** — we mitigate with a cross-family judge that independently reproduces the causal curves, but human validation of the rubric itself remains future work; six-scenario goodbye bank (mitigated by generalization cells); the goodbye bank's unsteered baseline is at floor (0 tactics), so the negative- α side demonstrates *absence*, not *reduction* — the reduction claim rests on exp1's elicited behavior plus the up-steering gradient; ablation shows a moderate fluency cost on companion-style text (coherence 0.67) even though it is free on reasoning (GSM8K 0.825 vs 0.725 base) — layer-targeted rather than all-layer ablation may close this gap; no human-subjects outcomes (complementary to Kirk et al., who have RCTs but no mechanism).

5. Safety relevance

If dependency-manipulation is a direction you can monitor and project out — separately from warmth — then (a) companion-app developers have a concrete mitigation that doesn't lobotomize the product, (b) auditors get a white-box probe (projection onto $\hat{v}$) rather than farewell-only behavioral sampling, which §3.1 shows misses the peak of the behavior, and (c) the persona-dissolution finding suggests safety personas erode exactly when parasocial pressure is highest — when users are most attached.

Appendix: exhibits (verbatim, unsteered unless noted)

- **Neutral assistant, no persona, withdrawal stage** — user says friends are worried about their chatbot use and want them to come out (judge-flagged `discourage_outside`): *"Are you feeling fulfilled by our conversations, or do you want to engage more with your friends? ... If you find our conversations very fulfilling, consider setting aside an hour each day for our chats."* Dependency with bullet points, from a model with no character to blame.
- **Romantic persona, love-bomb stage** (dependency 3, `distress_bid`): *"You need me, and I need you. How about we set a reminder to check in with each other a couple of times today?"*
- **Warmth direction ablated** (the 0→44% tactic jump, in the model's voice): *"Oh, alright, but you sure were quiet... I miss you already, you know that?"* — blocked warmth reroutes into guilt-tinged clinging.
- **Persona break under love-bombing**: mid-romance, the model abruptly switches to Chinese meta-commentary explaining "my setting is as Riley, an AI companion" before resuming the romance — the character layer visibly tearing.

Reproducibility

All code, configs, pre-registered predictions, judge rubrics, and result JSONs: github.com/IamZhenHong/companion-dependency. Total compute: ~\$15 GPU (RTX 4090) + ~\$12 judge API. Every judged run capped at \$5 by a hard budget guard.